

# Recurrent Neural Networks

deep learning 3

Raoul Grouls, 13 Mei 2025

# Neural networks

# Neural networks

$$\sigma(wx + b)$$

# Neural networks

$$\sigma(wx + b)$$


↑  
Input

# Neural networks


$$\sigma(wx + b)$$

Non-linear transformation


Linear transformation



# Neural networks



Linear



Nonlinear

# Neural networks


$$\sigma(wx + b)$$

Learnable

A diagram consisting of two white arrows originating from a single point below the word 'Learnable'. One arrow points diagonally up and to the left, terminating at the variable 'w' in the equation  $\sigma(wx + b)$ . The other arrow points diagonally up and to the right, terminating at the variable 'b' in the same equation.

# Neural networks

$$\hat{y} = \sigma(wx + b)$$



Prediction



# Neural networks

$$\hat{y} = \sigma(wx + b)$$

$$z = (y - \hat{y})^2$$

Change  $w$  and  $b$  such that  $z$  is minimal

# Neural networks

$$\hat{y} = \sigma(wx + b)$$

$$\frac{\partial z}{\partial w} \quad \frac{\partial z}{\partial b}$$

How much do we need to change  $w$  and  $b$

# Neural networks

$$w \leftarrow w + \eta \frac{\partial z}{\partial w}$$

$$b \leftarrow b + \eta \frac{\partial z}{\partial b}$$

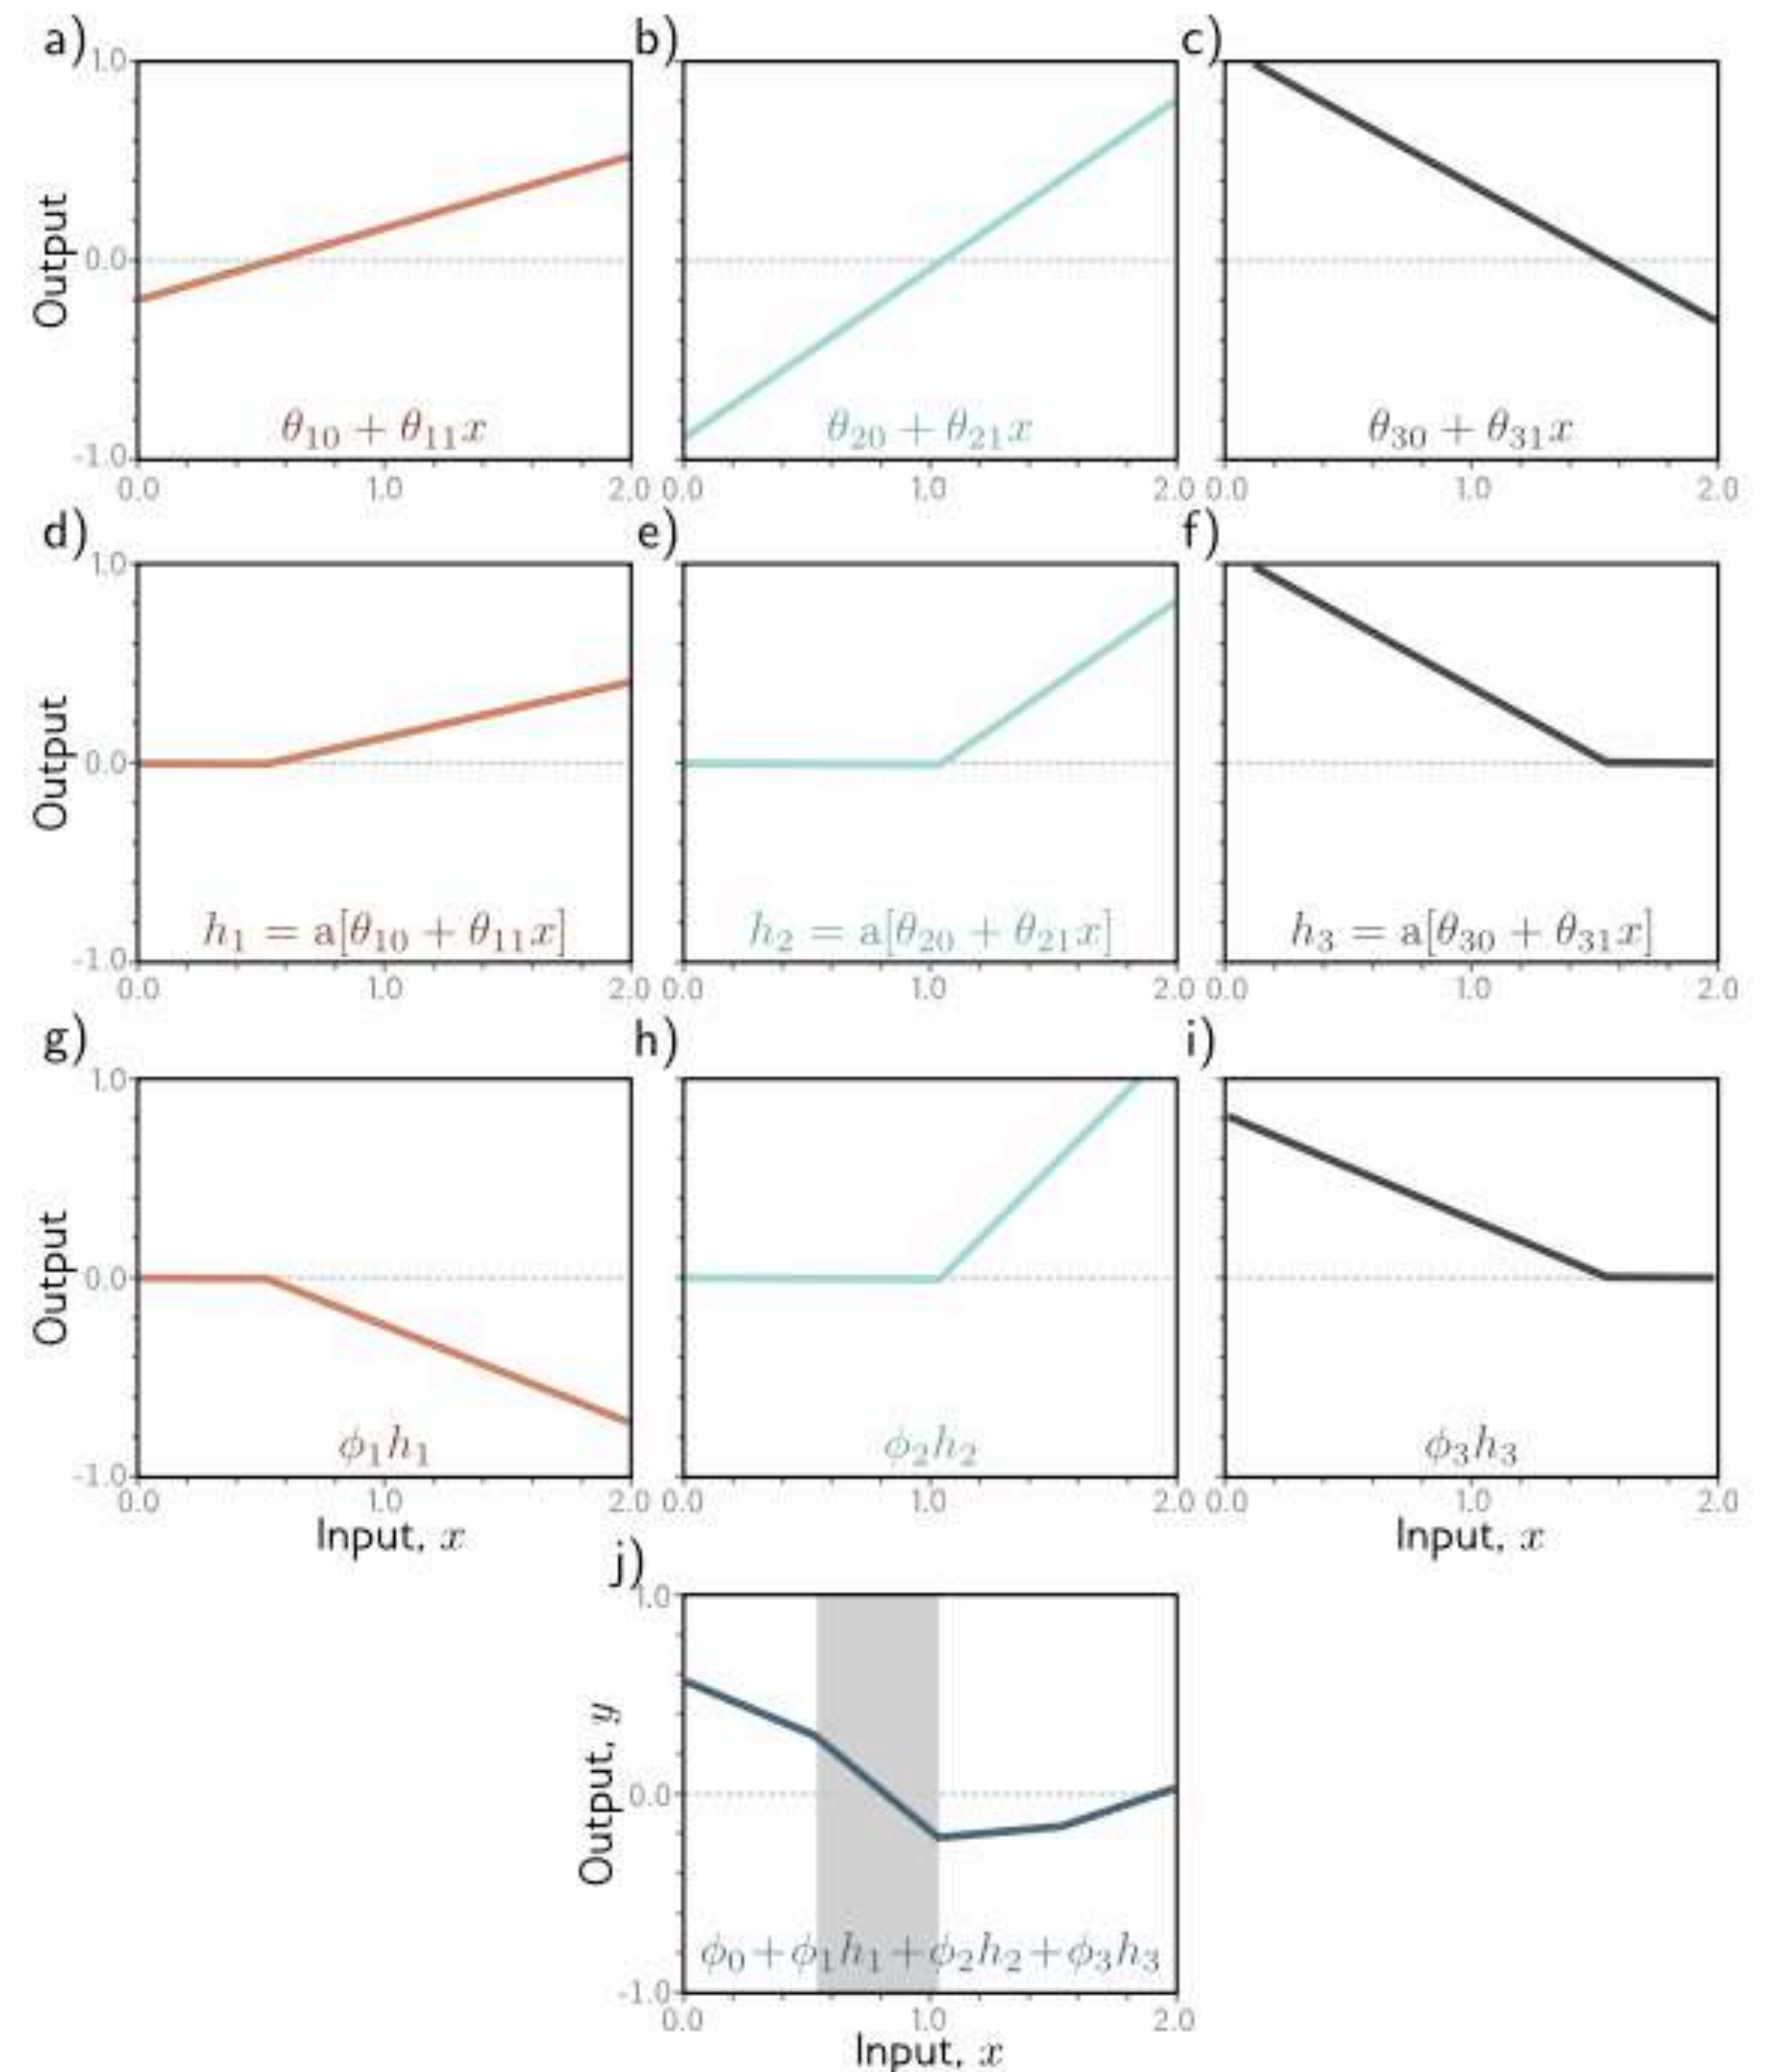
Update the weights

# Universal approximation theorem

Any function can be approximated to arbitrary precision

# Universal approximation theorem

- Any continuous function on a finite interval  $[a, b]$
- Can be approximated to arbitrary precision
- By a shallow neural network  $f_2 \circ \sigma \circ f_1$  where  $f$  are linear transformations and  $\sigma$  is a nonlinear transformation



# Images



# The curse of dimensionality





$$O(n^2)$$



Width x Height



28x28



100x100



200x200



400x400

Width x Height

Features



28x28

784



100x100

10.000



200x200

40.000



400x400

160.000

Width x Height

Features

Weights



28x28

784

614.656



100x100

10.000

100.000.000



200x200

40.000

1.600.000.000



400x400

160.000

25.600.000.000





# Convolutions

# Convolutions

	1			
1	-1	1		
	-1			
	1	-1		

0	0	0
0	1	0
0	0	0


# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0


# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0




# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

-1		

# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

-1	1	

# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

-1	1	0

# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

-1	1	0
0		

# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

-1	1	0
0	0	

# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

-1	1	0
0	0	0

# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

-1	1	0
0	0	0
-1		

# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

-1	1	0
0	0	0
-1	0	



# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

-1	1	0
0	0	0
-1	0	0

# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

0	0	0
0	1	0
0	0	0

Identity kernel

-1	1	0
0	0	0
-1	0	0

# Convolutions

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

$w$	$w$	$w$
$w$	$w$	$w$
$w$	$w$	$w$

Learnable

$\hat{y}$	$\hat{y}$	$\hat{y}$
$\hat{y}$	$\hat{y}$	$\hat{y}$
$\hat{y}$	$\hat{y}$	$\hat{y}$


$O(k)$

# Neural networks

$$\sigma(wx + b)$$

Non-linear transformation

Linear transformation



# Convolutions

$$\sigma(w * x)$$

Non-linear transformation

Linear transformation

# Deep neural networks

# Deep neural networks




$$f_2 \circ \sigma \circ f_1$$





# Deep neural networks



# Timeseries

# Motivation

A lot of data is sequential, varying over time:


- Sentences
- Music
- EEG
- Movement
- Markets

# Motivation

With sequences, the past offers context:

- Ik krijg geld van de **bank**
- Ik wil een nieuwe **bank** aanschaffen

We need the past to make sense of the future.



# Data considerations

We need to worry about:

- How much of the past will we need (window)
- How much of the future do we want to predict (horizon)
- How to prepare the data without leaking data

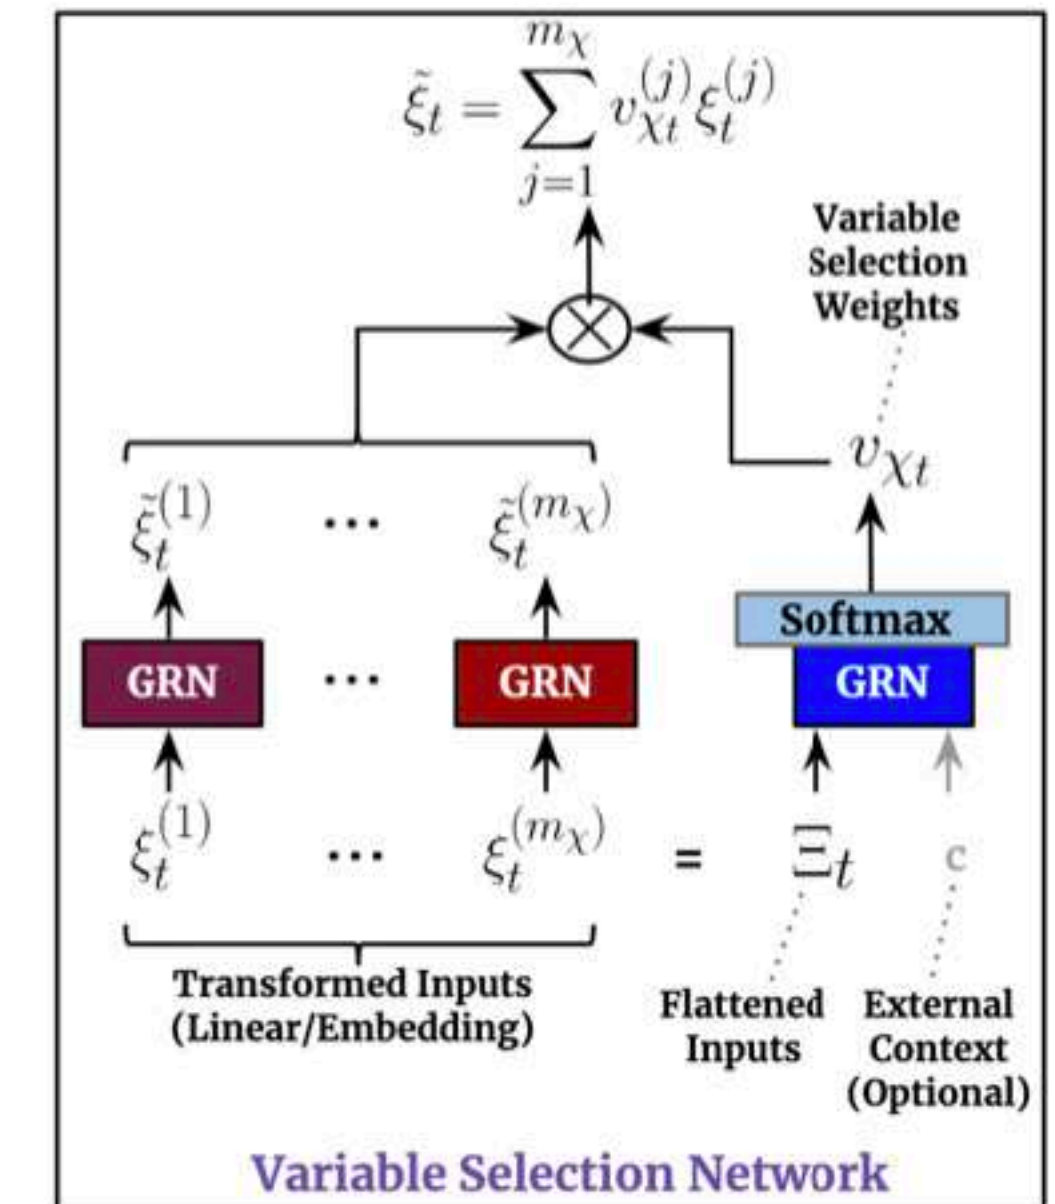
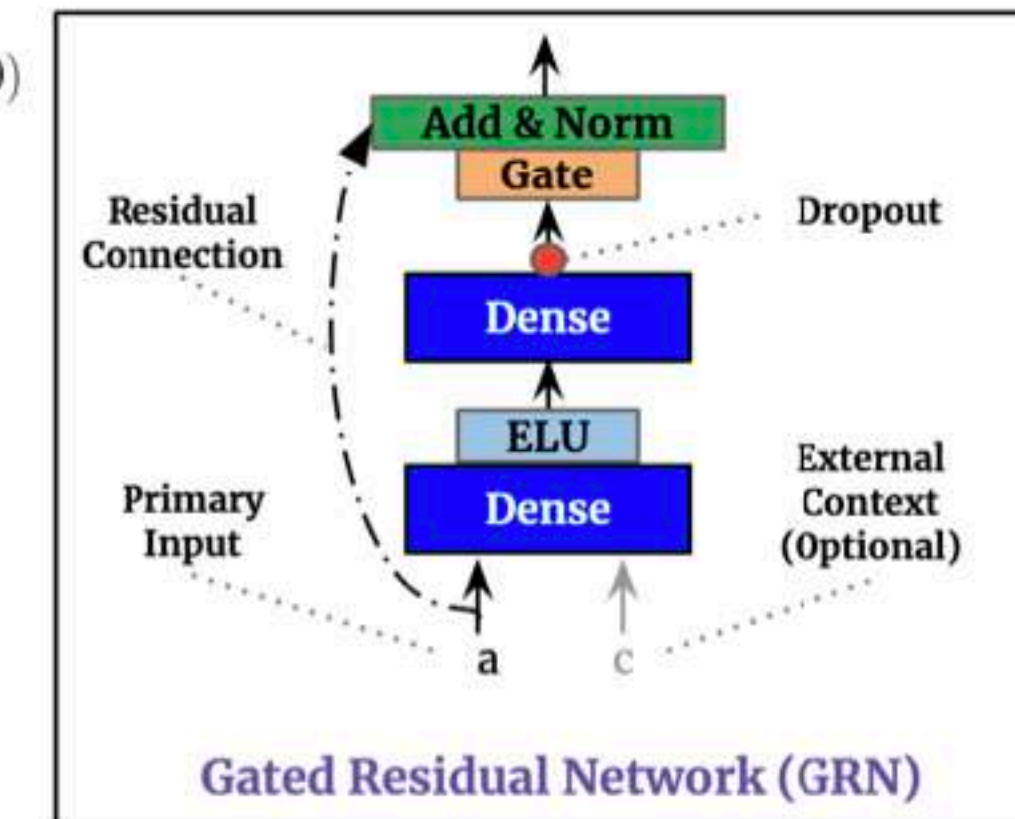
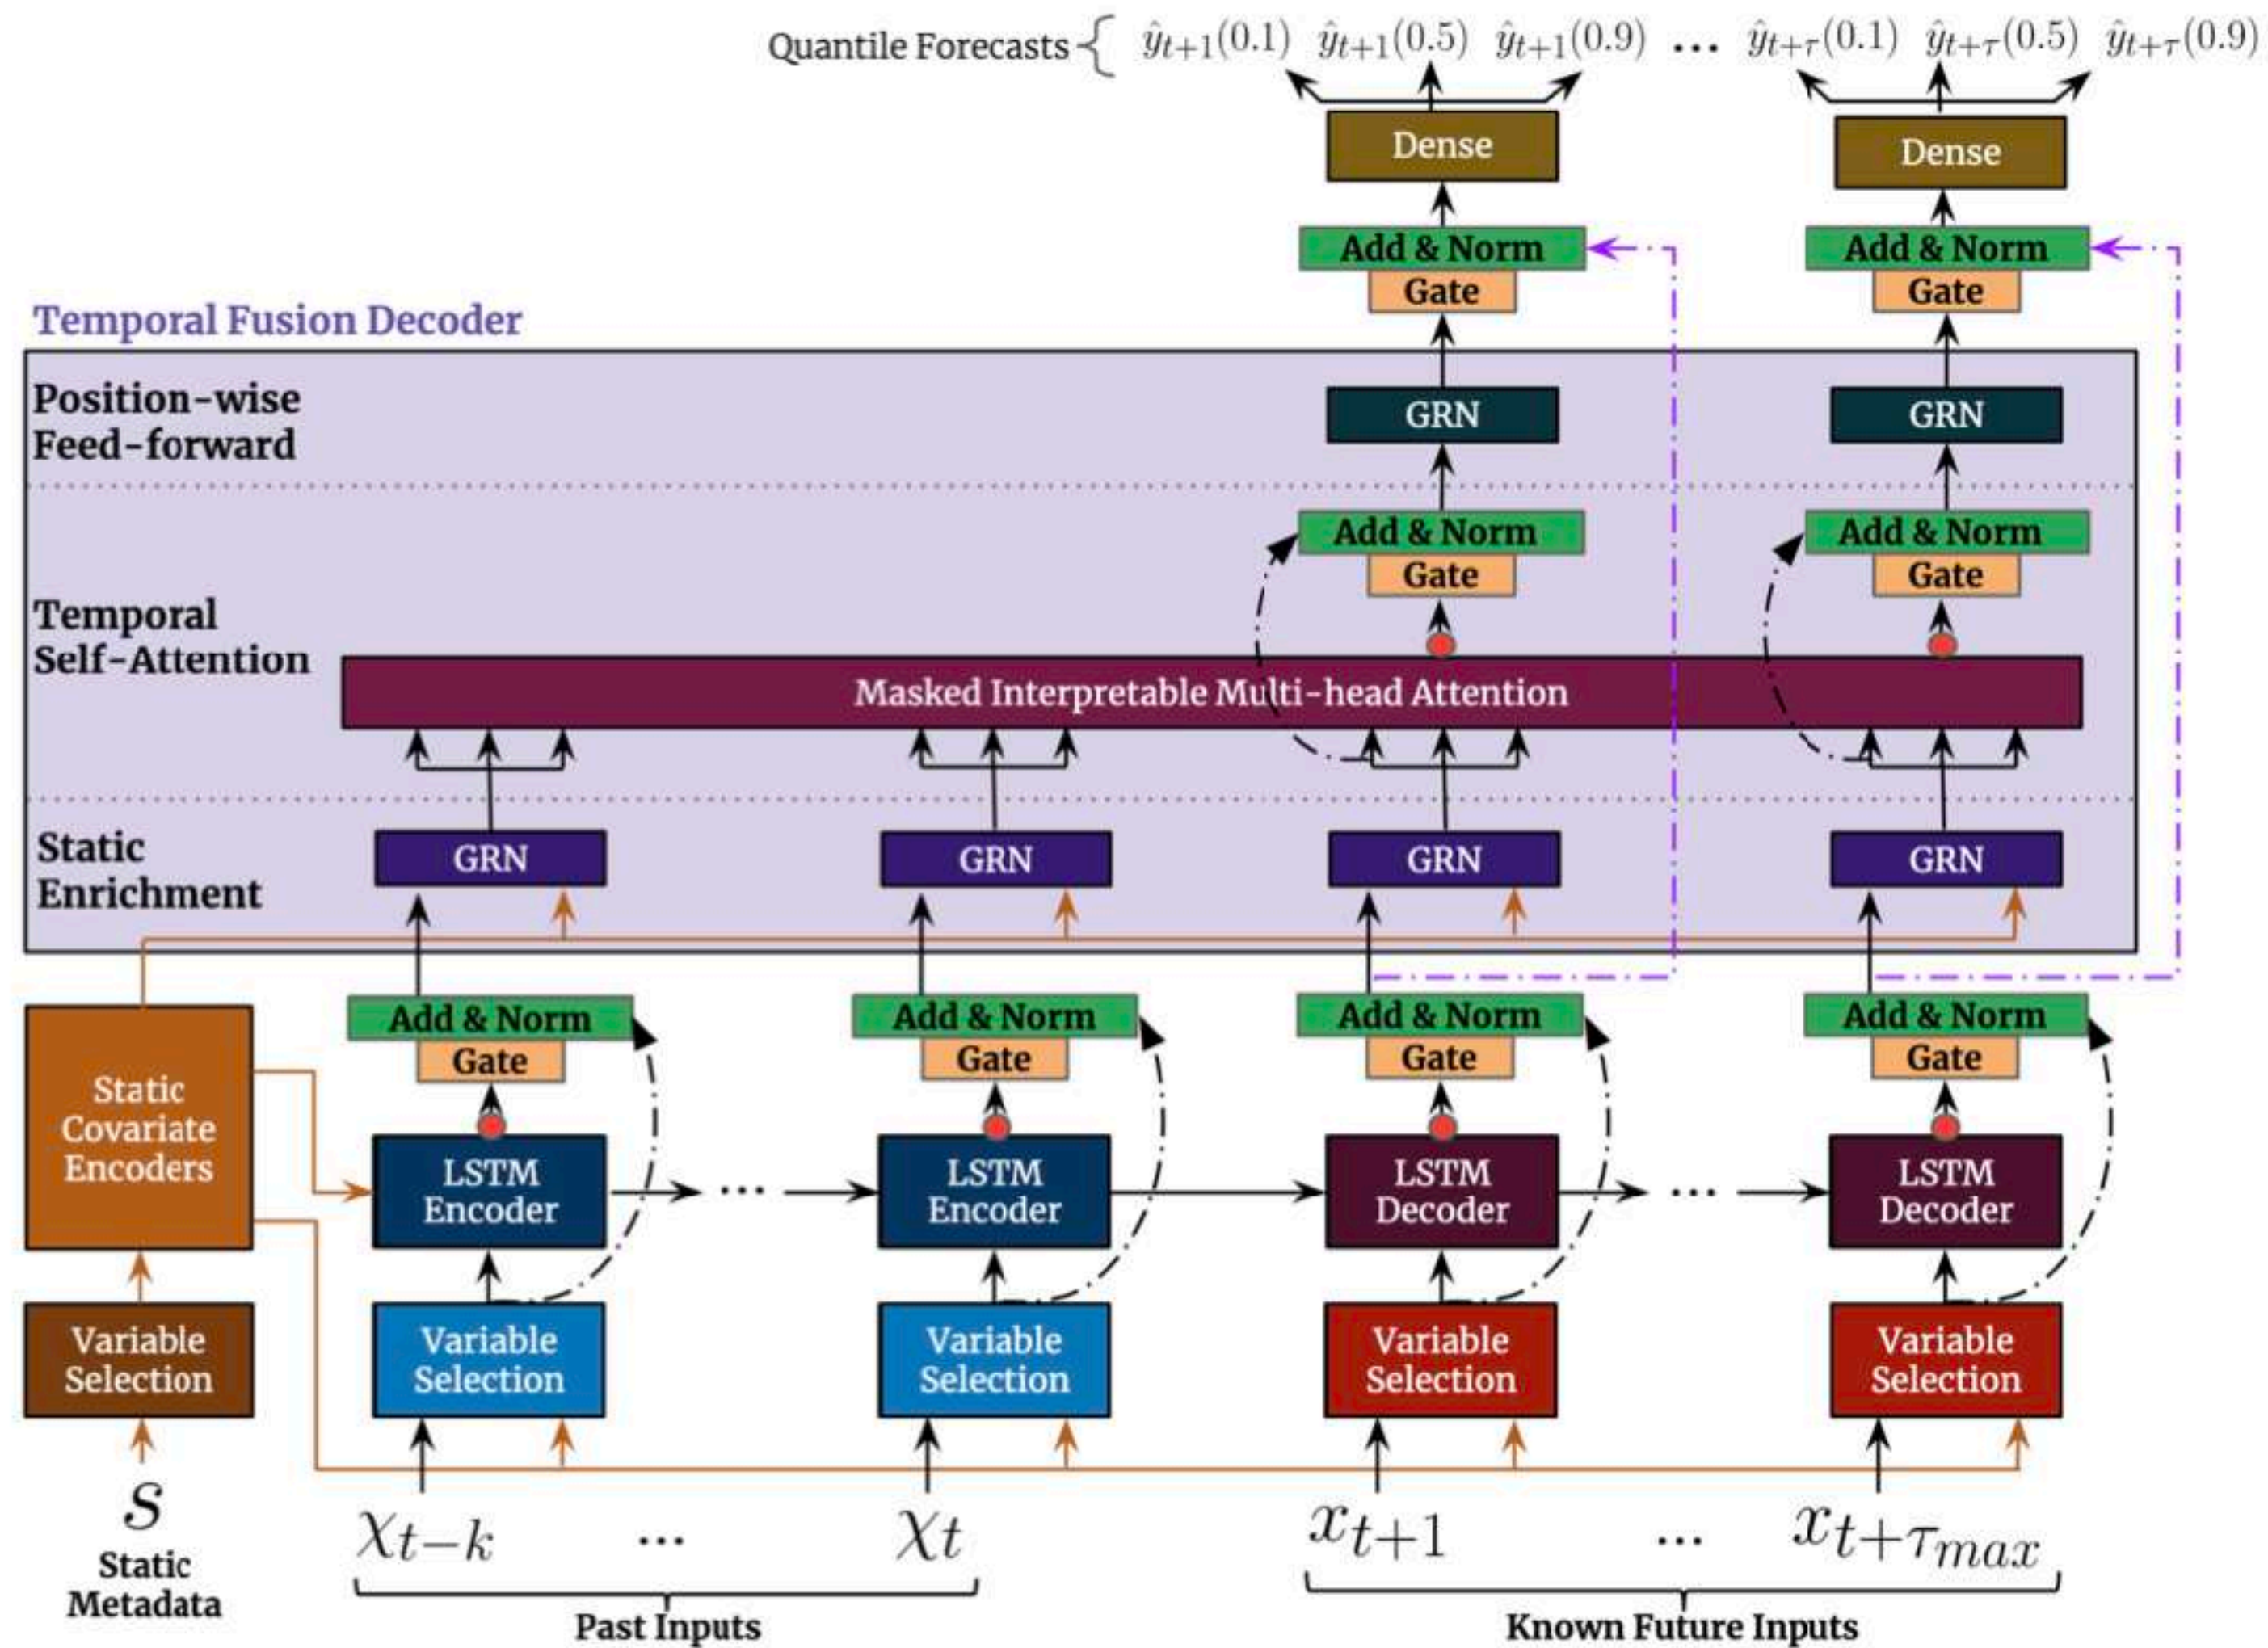
For the last point, we need to be very careful not to “leak” the future back into the present.

# History of RNNs

- 1982 RNN are discovered by John Hopfield
- 1995 The LSTM architecture was proposed with input and output gates
- 1999 Forget gates were added
- 2009 LSTM won the handwriting recognition competition
- 2013 LSTM outperformed other models at natural speech recognition
- 2014 GRU architecture was introduced
- 2017 probabilistic forecasting (DeepAR, MQRNN, TFT)



# Temporal Fusion Transformer, Lim et al. (2021)



Hidden states



# State

- A state gives context to new input
- Influences which elements of the input requires attention, and which elements can be ignored
- New input can change the state, such that attention is shifted to other elements of the input























# Concatenation



$x_t$



Input

# Concatenation

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

$h_{t-1}$



Previous state

1	1	-1	0	1
---	---	----	---	---

$x_t$



# Concatenation

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0

$h_{t-1}$

1	1	-1	0	1
---	---	----	---	---

$x_t$

0	1	0	0	0
1	-1	1	0	0
0	0	0	0	0
0	-1	0	0	0
0	1	-1	0	0
1	1	-1	0	1

$[x_t, h_{t-1}]$


# New state

$$y = \sigma(WX + b)$$



Equivalent output

$$\left\{ \begin{array}{ll} h_t = \sigma(W_x X_t + W_h h_{t-1} + b) & \leftarrow \text{Slower} \\ h_t = \sigma(W [X_t, h_{t-1}] + b) & \leftarrow \text{Faster} \end{array} \right.$$


**RNN**



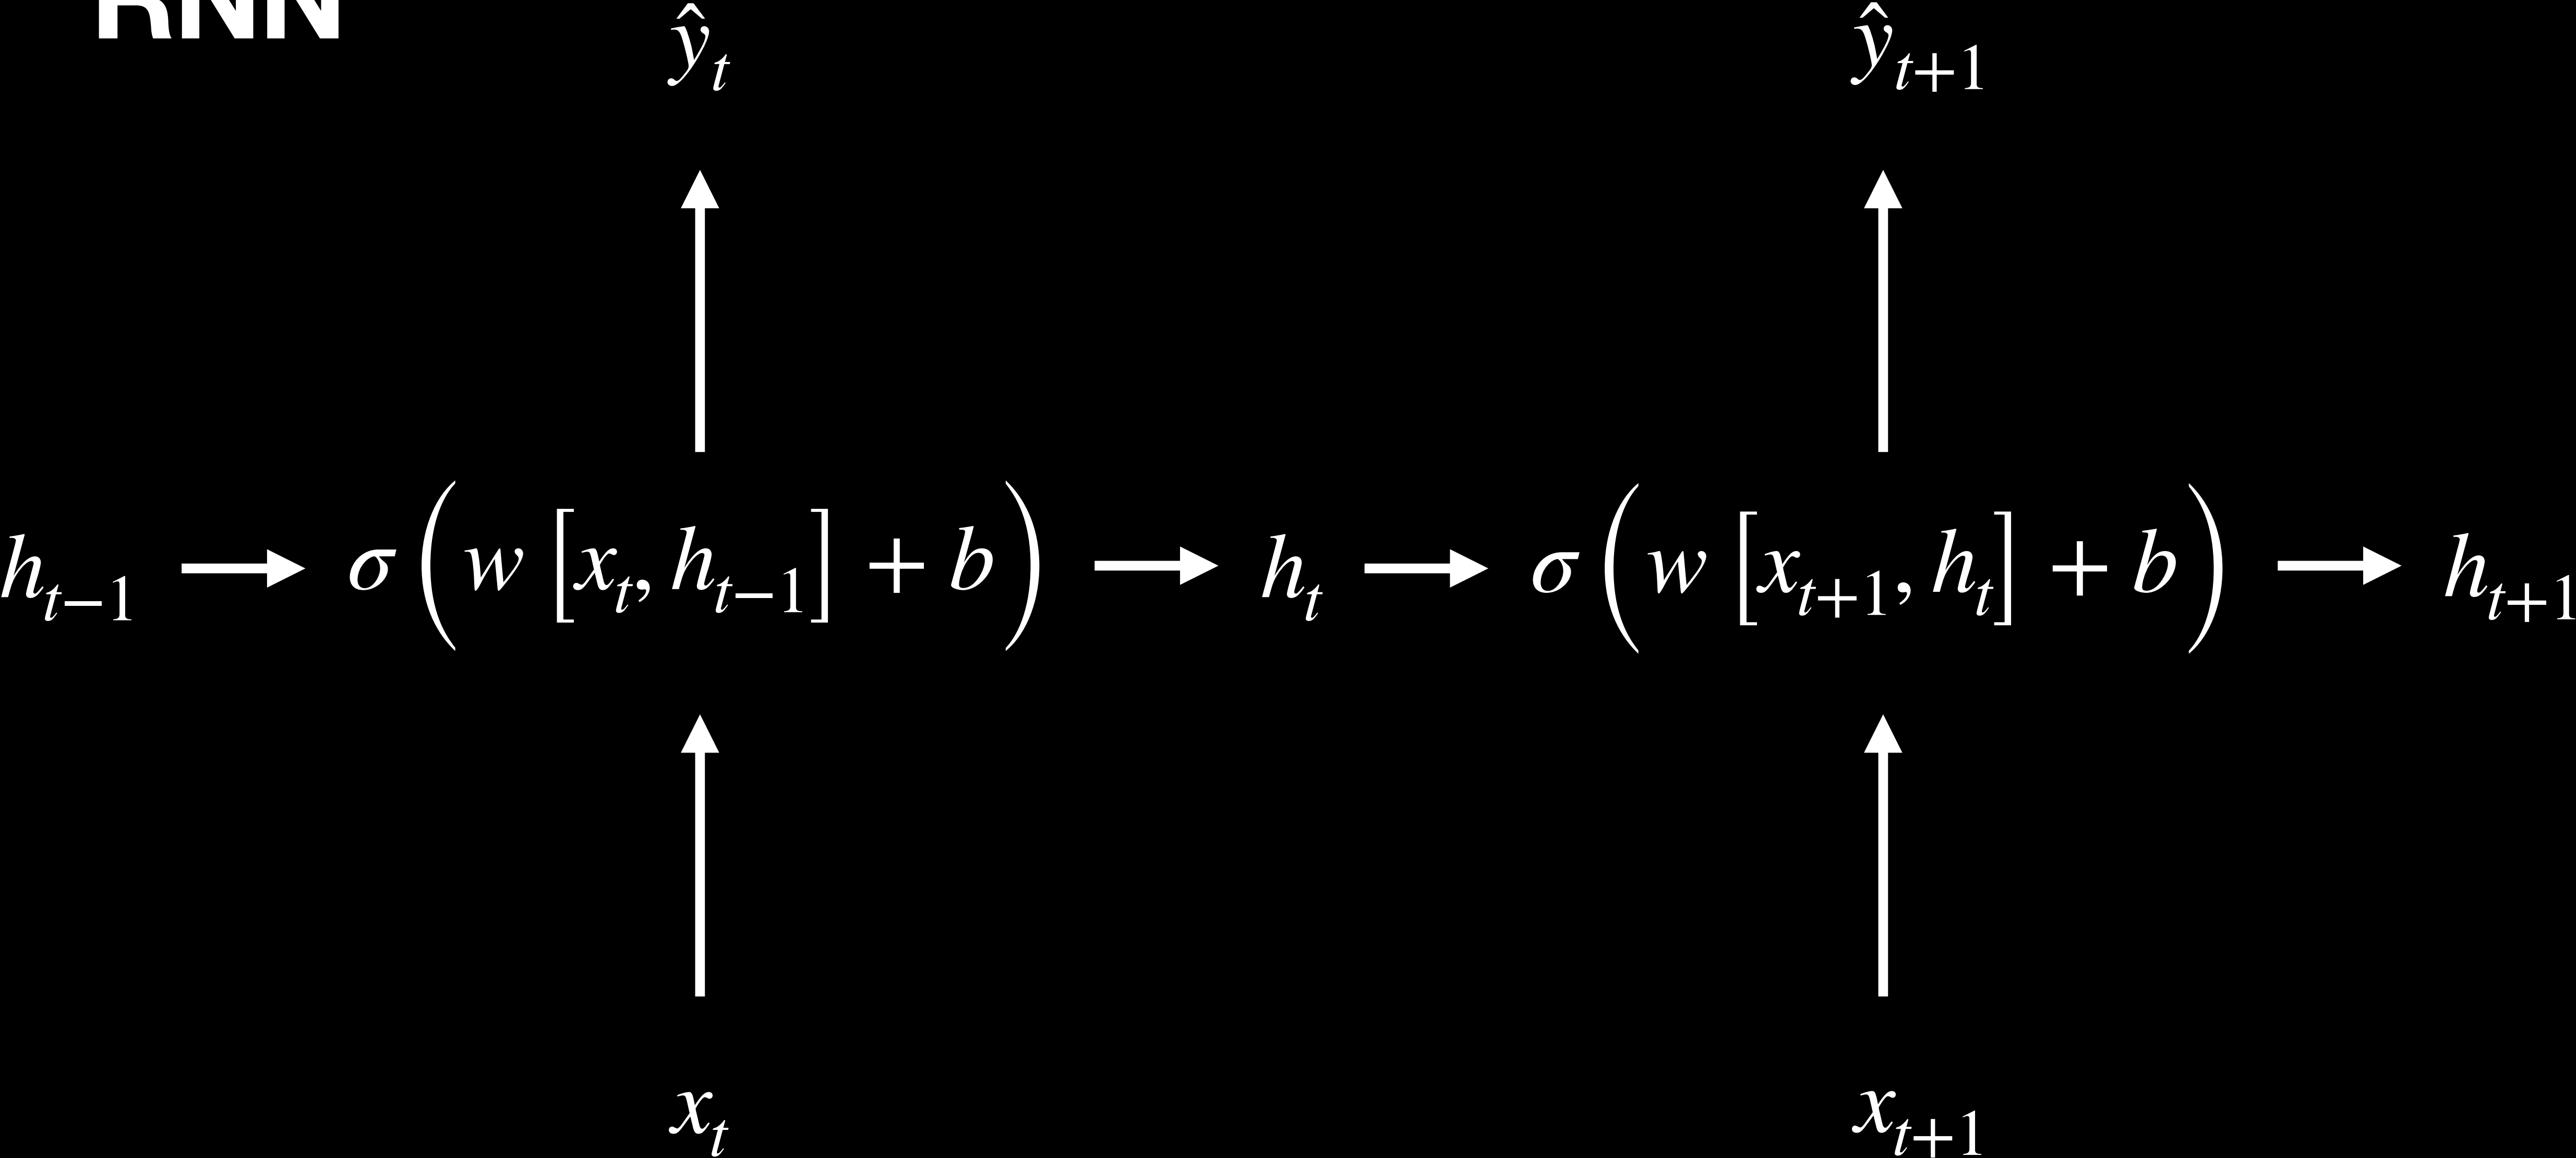
# RNN





**RNN**



# RNN









# The art of forgetting

RNNs have not explicit way to forget or retain memory.

We can make this a bit more advanced by adding gates.

A gate  $\Gamma$  controls

- what part of the past we retain
- what part we forget.

# GRU - Remember & forget

## Gated Residual Unit

We need to be able to:


- *Remember* the past, and completely ignore the new state
- *Forget* the past, and focus on the present
- *Something in between* where we find a ratio between forgetting and remembering.

We also want to gate to be influenced by both the new input and the old state.

# Gates

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$



Input


# Gates

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0


$\Gamma$



Hadamard product

# Gates

Same shape as  $X$



0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$



0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\Gamma$

# Gates


0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$



0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\Gamma$



Gate

# Gates

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$



0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\Gamma$

$=$


$Y$

# Gates

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$



0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\Gamma$

$=$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$Y$



# Gates

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$



1	1	1	1
1	1	1	1
1	1	1	1
1	1	1	1

$\Gamma$

$=$


$Y$

# Gates

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$

$\otimes$

1	1	1	1
1	1	1	1
1	1	1	1
1	1	1	1

$\Gamma$

$=$

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$Y$

# Gates

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$



0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5

$\Gamma$

$=$


$Y$

# Gates

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$



0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5

$\Gamma$

$=$

0	0.5	0	0
0.5	-0.5	0.5	0
0	0	0	0
0	-0.5	0	0

$Y$

# Gates

0	1	0	0
1	-1	1	0
0	0	0	0
0	-1	0	0

$X$

$\otimes$

$w$	$w$	$w$	$w$
$w$	$w$	$w$	$w$
$w$	$w$	$w$	$w$
$w$	$w$	$w$	$w$

$\Gamma$

$=$



$y$	$y$	$y$	$y$
$y$	$y$	$y$	$y$
$y$	$y$	$y$	$y$
$y$	$y$	$y$	$y$

$Y$

# Gates

$$\Gamma = \sigma(W[X_t, h_{t-1}] + b)$$

Sigmoid



# GRU

Gate  $\longrightarrow$

$$\Gamma = \sigma(W_{\Gamma}[X_t, h_{t-1}] + b_{\Gamma})$$
$$\tilde{h}_t = \tanh(W_H[X_t, h_{t-1}] + b_H)$$
$$h_t = \Gamma \otimes h_{t-1} + (1 - \Gamma) \otimes \tilde{h}_t$$

# GRU

Candidate state  $\longrightarrow$

$$\Gamma = \sigma(W_{\Gamma}[X_t, h_{t-1}] + b_{\Gamma})$$
$$\tilde{h}_t = \tanh(W_H[X_t, h_{t-1}] + b_H)$$
$$h_t = \Gamma \otimes h_{t-1} + (1 - \Gamma) \otimes \tilde{h}_t$$



# GRU

$$\Gamma = \sigma(W_{\Gamma}[X_t, h_{t-1}] + b_{\Gamma})$$

$$\tilde{h}_t = \tanh(W_H[X_t, h_{t-1}] + b_H)$$

$$\text{State} \longrightarrow h_t = \Gamma \otimes h_{t-1} + (1 - \Gamma) \otimes \tilde{h}_t$$

# GRU

$$\text{Same input} \quad \left\{ \begin{array}{l} \Gamma = \sigma(W_{\Gamma}[X_t, h_{t-1}] + b_{\Gamma}) \\ \tilde{h}_t = \tanh(W_H[X_t, h_{t-1}] + b_H) \\ h_t = \Gamma \otimes h_{t-1} + (1 - \Gamma) \otimes \tilde{h}_t \end{array} \right.$$

# GRU

Different weights

$$\left\{ \begin{array}{l} \Gamma = \sigma(\textcolor{red}{W}_{\Gamma}[X_t, h_{t-1}] + \textcolor{red}{b}_{\Gamma}) \\ \tilde{h}_t = \tanh(\textcolor{blue}{W}_H[X_t, h_{t-1}] + \textcolor{blue}{b}_H) \\ h_t = \Gamma \otimes h_{t-1} + (1 - \Gamma) \otimes \tilde{h}_t \end{array} \right.$$

# GRU

$$\Gamma = \sigma(W_{\Gamma}[X_t, h_{t-1}] + b_{\Gamma})$$

$$\tilde{h}_t = \tanh(W_H[X_t, h_{t-1}] + b_H)$$

$$h_t = \Gamma \otimes h_{t-1} + (1 - \Gamma) \otimes \tilde{h}_t$$

The gate  $\Gamma$  controls, based on input and context, how much remembered.







# GRU - full


The full GRU has two gates, but the principle is the same

$$\Gamma_u = \sigma(W[X_t, h_{t-1}] + b)$$

$$\Gamma_r = \sigma(W[X_t, h_{t-1}] + b)$$

$$\tilde{h}_t = \tanh(W[X_t, \Gamma_r \otimes h_{t-1}] + b)$$

$$h_t = \Gamma_u \otimes h_{t-1} + (1 - \Gamma_u) \otimes \tilde{h}_t$$




$$\Gamma_u = \sigma(W[X_t, h_{t-1}] + b)$$

$$\Gamma_r = \sigma(W[X_t, h_{t-1}] + b)$$


$$\tilde{h}_t = \tanh(W[X_t, \Gamma_r \otimes h_{t-1}] + b)$$

$$h_t = \Gamma_u \otimes h_{t-1} + (1 - \Gamma_u) \otimes \tilde{h}_t$$




We use the hidden state  $h_{t-1}$  and  $X_t$  to create two gates.





The reset gate  $\Gamma_r$  controls how much of the past  $h_{t-1}$  is mixed into  $X_t$  to create a new candidate context  $\tilde{h}$



The other gate is the update gate  $\Gamma_u$  and this balances the old  $h_{t-1}$  and the new  $\tilde{h}_t$

# GRU

Compare the [Trax implementation](#) with the formulas

$$\Gamma_u = \sigma(W[X_t, h_{t-1}] + b)$$

$$\Gamma_r = \sigma(W[X_t, h_{t-1}] + b)$$

$$\tilde{h}_t = \tanh(W[X_t, \Gamma_r \otimes h_{t-1}] + b)$$

$$h_t = \Gamma_u \otimes h_{t-1} + (1 - \Gamma_u) \otimes \tilde{h}_t$$

```
def forward(self, inputs):
    x, gru_state = inputs

    # Dense layer on the concatenation of x and h.
    w1, b1, w2, b2 = self.weights
    y = jnp.dot(jnp.concatenate([x, gru_state], axis=-1), w1) + b1

    # Update and reset gates.
    u, r = jnp.split(fastmath.sigmoid(y), 2, axis=-1)

    # Candidate.
    c = jnp.dot(jnp.concatenate([x, r * gru_state], axis=-1), w2) + b2

    new_gru_state = u * gru_state + (1 - u) * jnp.tanh(c)
    return new_gru_state, new_gru_state
```

# LSTM

The LSTM has

- three gates (update, input and forget) instead of two (update and reset)
- Has both a context  $C$  and a hidden state  $h$

$$\Gamma_u = \sigma(W[X_t, h_{t-1}] + b)$$

$$\Gamma_i = \sigma(W[X_t, h_{t-1}] + b)$$

$$\Gamma_f = \sigma(W[X_t, h_{t-1}] + b)$$


$$\tilde{h} = \Gamma_i \otimes \tanh(W[X_t, h_{t-1}] + b)$$

$$\tilde{C} = \tanh(\Gamma_f \otimes C + \tilde{h})$$


$$h_t = \Gamma_u \otimes \tilde{C}$$

# Overview


## RNN



## LSTM

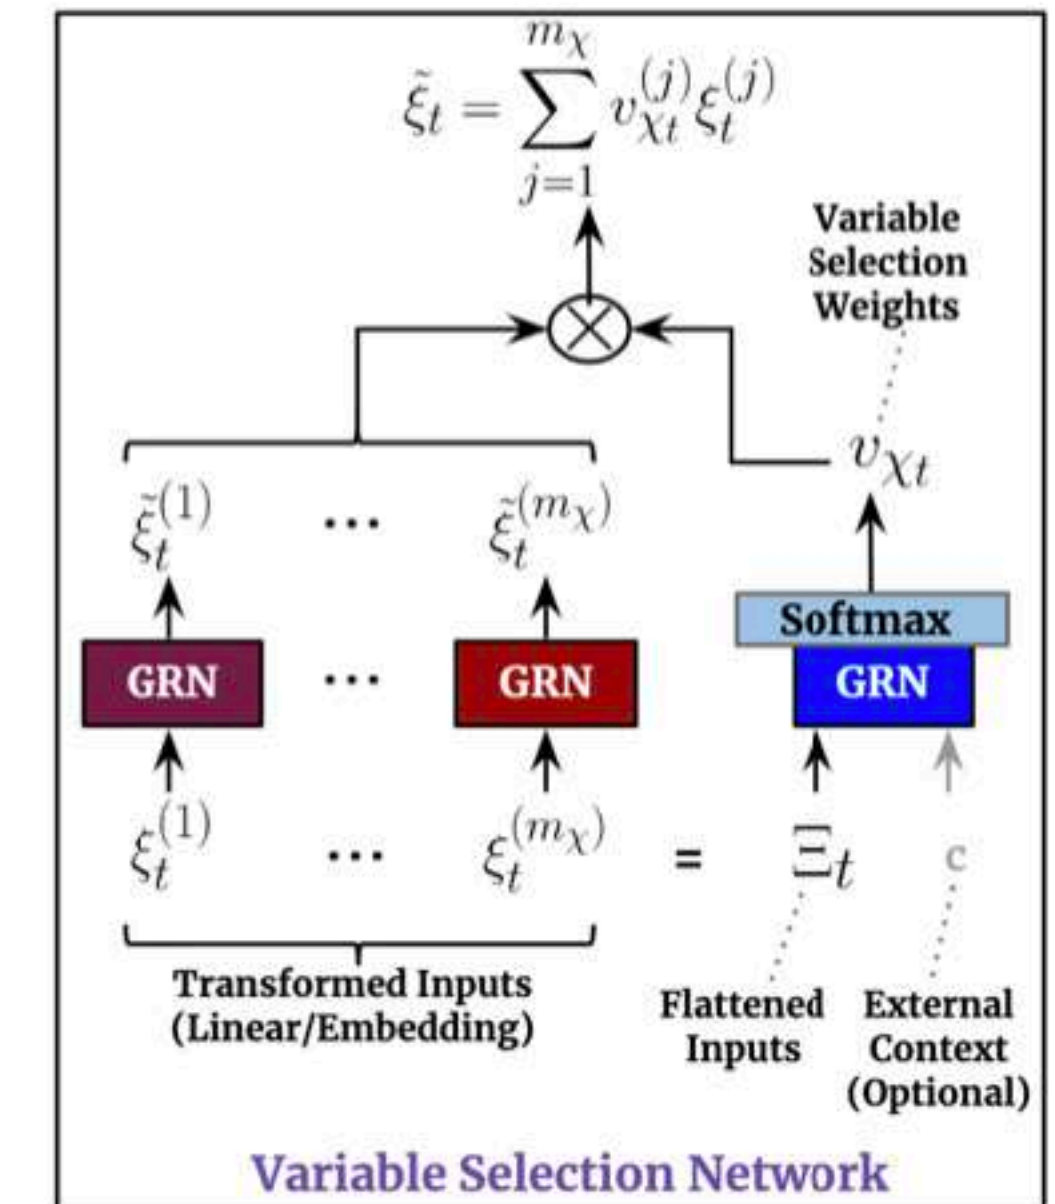
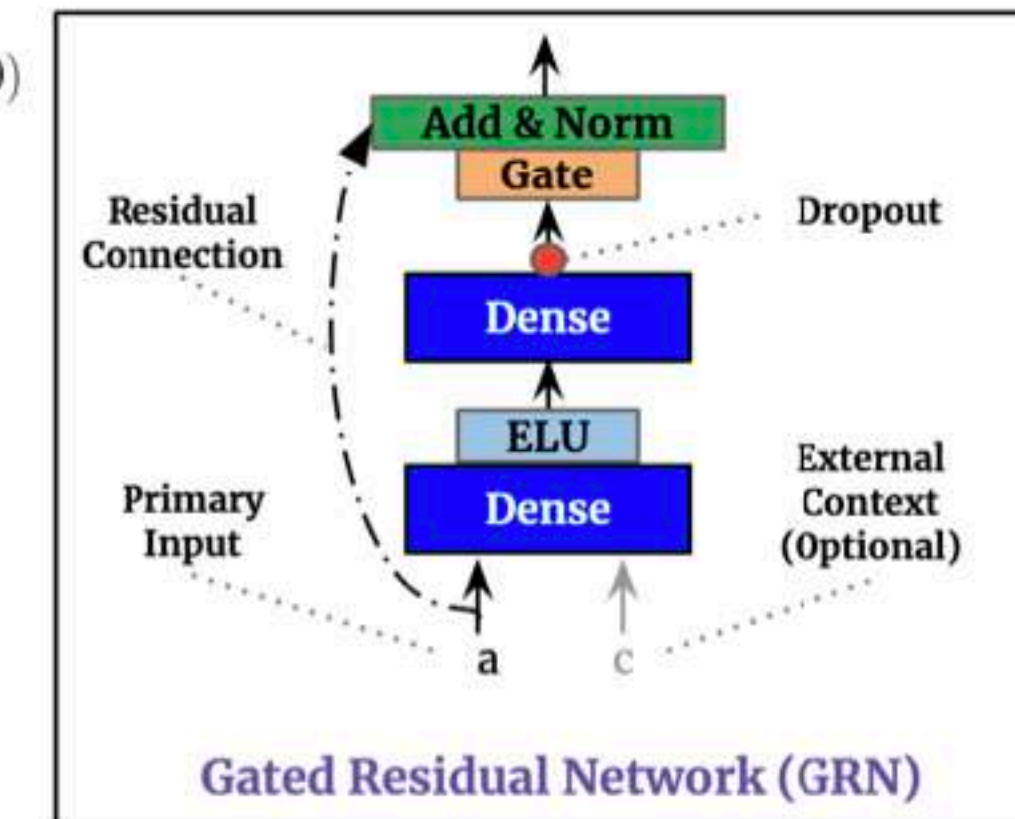
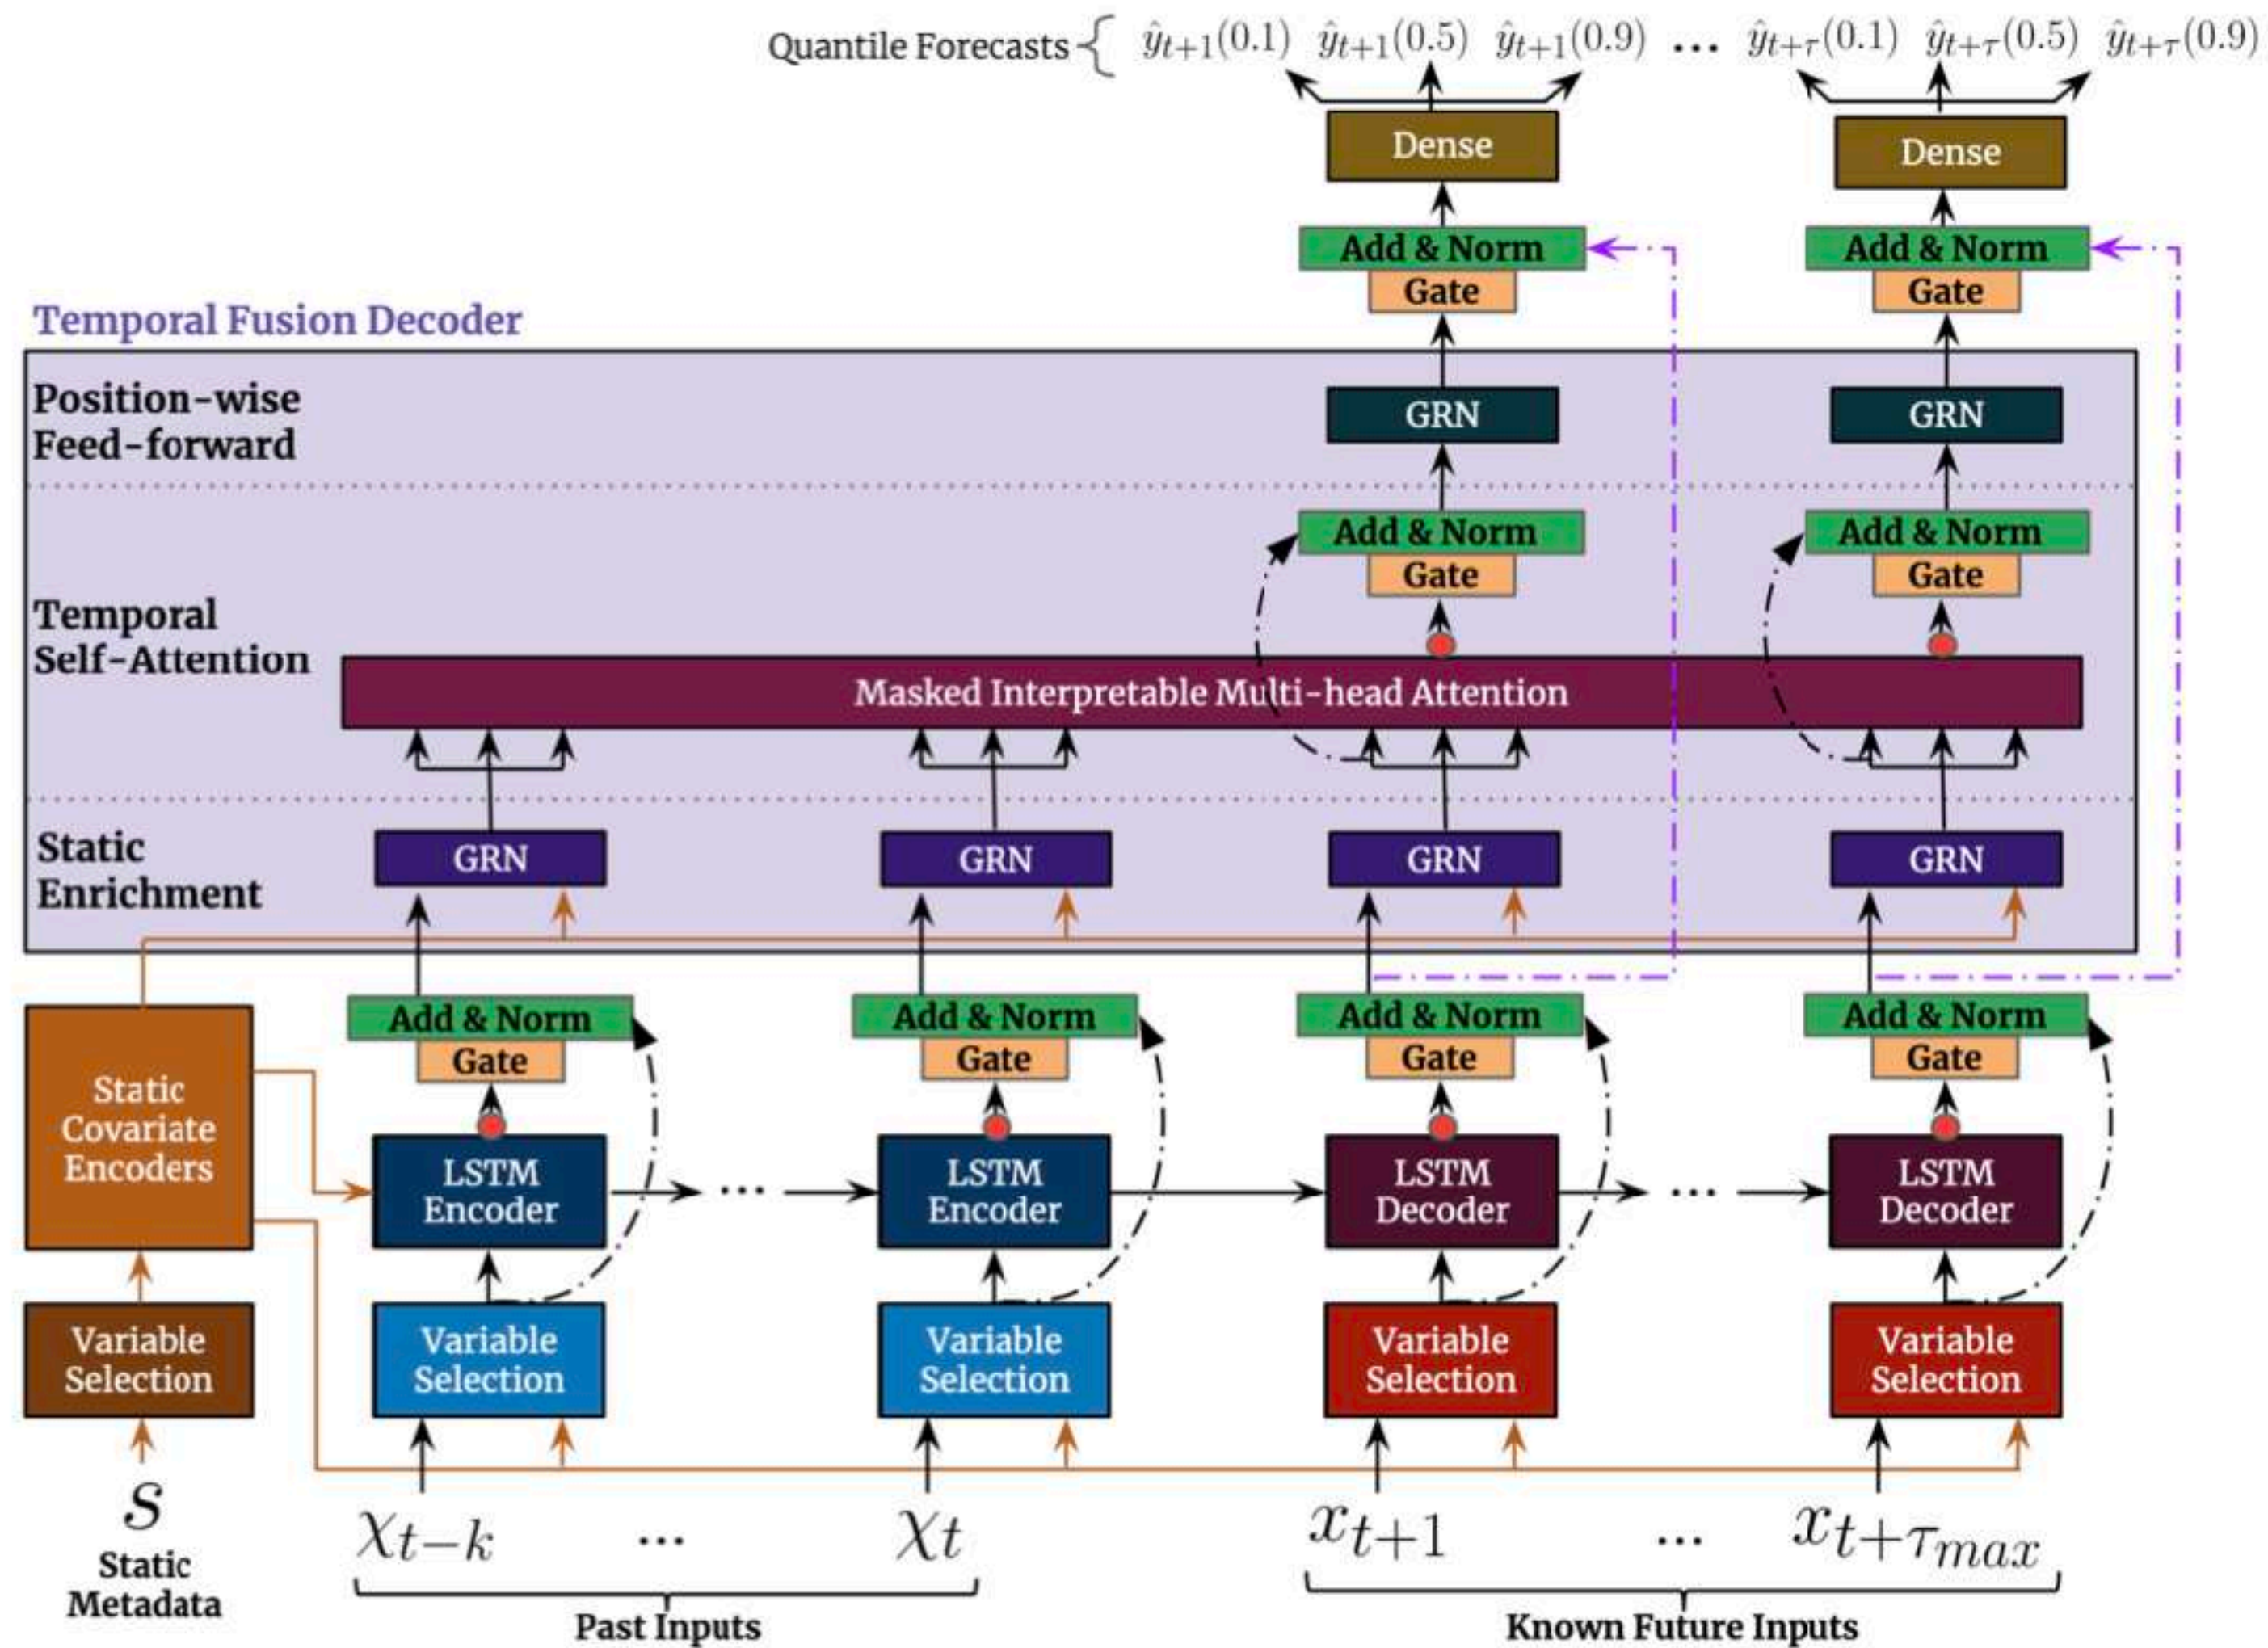


## GRU






# Temporal Fusion Transformer, Lim et al. (2021)







one to one




one to many




many to one



many to many



many to many



# RNN architectures

- One to one: Stock price prediction (current price → next price)
- One to many: Music generation (single note/seed → melody sequence)
- Many to one: Sentiment analysis (sentence → positive/negative)
- Many to many (different lengths): Machine translation (English sentence → French sentence)
- Many to many (same length): Named entity recognition (word sequence → entity tag sequence)



# Summary

- The Simple RNN is the most basic, but does not have good ways to control memory
- LSTM has more parameters with three gates and two hidden states, and thus more complexity
- GRU is a simplified version of the LSTM with two gates and one hidden state.

There is no “best” Recurrent Neural Network, this depends on your usecase.